

CS/ECE 5381/7381
Computer Architecture
Spring 2023

Dr. Manikas

Computer Science

Lecture 4: Jan. 26, 2023

LOCKDOWN BROWSER

- Notes from this week's announcement
- Our exams will use Lockdown Browser
- Exam 1 will be in three weeks (**Feb. 16**)
- Please install and test Lockdown Browser on your computer
- Use Practice Exam to verify set up and to get familiar with Lockdown Browser

Assignments

- Quiz 2 – due Sat., Jan. 28 (11:59 pm)
 - Covers concepts from Module 2 (this week)

Quiz 2 Details

- The quiz is open book and open notes.
- You are allowed 90 minutes to take this quiz.
- You are allowed 2 attempts to take this quiz - your highest score will be kept.
 - Note that some questions (e.g., fill in the blank) will need to be graded manually
- Quiz answers will be made available 24 hours after the quiz due date.

Instruction Set Principles

(Appendix A, Hennessy and Patterson)

Note: some course slides adopted from
publisher-provided material

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Memory Addressing

- For any ISA class, must define how memory is addressed
 - Address interpretation
 - Address specification

Interpreting Memory Addresses

- Memory size
 - $2^{10} = 1 \text{ K (kilo)}$
 - $2^{20} = 1 \text{ M (mega)}$
 - $2^{30} = 1 \text{ G (giga)}$
 - $2^{40} = 1 \text{ T (tera)}$
 - $2^{50} = 1 \text{ P (peta)}$
- $1 \text{ B (byte)} = 8 \text{ b (bits)}$

Word Size

- CPU word size = w
 - Typical: $w = 16$ or 32 bits
- Memory word size = s
 - Typical: $s = 8$ bits (memory is *byte-addressable*)
- Need to identify order of memory words in CPU word

Word-Ordering Schemes

- Big Endian
 - Stores the **MSB** ("the big end") at the **lower** byte address and the **LSB** at the **higher** byte address
- Little Endian
 - Stores the **LSB** ("the little end") at the **lower** byte address and the **MSB** at the **higher** byte address
- Example:
 - store ABCD1234H at word address location 0
 - AB is MSB, 34 is LSB

Memory Ordering Example

Big Endian

Byte Address	0	1	2	3
Data	AB	CD	12	34

Little Endian

Byte Address	0	1	2	3
Data	34	12	CD	AB

Addressing Modes

- What object is being addressed?
 - Memory location?
 - Specified by *effective address*
 - Register?
 - Specified by a *register number*
 - Constant?
 - Constant value contained in instruction

Register Mode

- All values in registers
- ADD R0, R1
 - $[R_x] = \text{contents of register } x$
 - $[R0] \leftarrow [R0] + R[1]$

Immediate Mode

- For constants
- `ADD R0, #3`
 - $[R0] \leftarrow [R0] + 3$

Register Indirect Mode

- Use a register value as a *pointer* to memory
- ADD R0, (R1)
 - $[R0] \leftarrow [R0] + \text{Mem}[[R1]]$
 - Where $\text{Mem}[x]$ is memory location x
 - Register R1 contains the value for x

Indexed Mode

- Add a value (index) to the memory address
- ADD R0, (R1+R2)
 - $[R0] \leftarrow [R0] + \text{Mem}[[R1] + [R2]]$
- Used in array addressing
 - R1 contains base array value
 - R2 contains index value

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Type and Size of Operands

- Character (8 bits)
- Half word (16 bits)
- Word (32 bits)
- Single-precision floating point (32 bits)
- Double-precision floating point (64 bits)

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Operations in the Instruction Set

- Data Transfer (load/store)
- Arithmetic/Logic
 - Add, subtract, multiply, divide, negation (2's comp)
 - And, or, not (1's comp)
- Control flow (branch, jump)
- Other
 - System, graphics, etc.

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Instructions for Control Flow

- Default instruction execution is *sequential*
- Control flow instructions allow instruction execution to go out of sequence
 - Conditional branches
 - Jumps (unconditional branches)
 - Procedure calls and returns

Control Flow Addressing

- Destination address of next instruction?
 - PC-relative: displacement added to the program counter value
 - Dynamic address: register (or other location) contains destination address

Conditional Branch Options

- Branch based on condition code (CC)
 - Special test bits (zero/nonzero, positive/negative)
- Branch based on value in condition register
- Compare and branch
 - CJNE – compare and jump if not equal

```
                LOAD R0, 1          ; initialize loop counter
LOOP:           (loop code)
                ADD R0,1             ;increment loop counter
                CJNE R0, 10, LOOP    ;repeat if loop counter < 10
                (next code)
```


Procedure Invocation Options

- Procedures: functions, subroutines
- Procedure Call
 - Save desired values from calling program
 - Save current PC value (so procedure knows where to return)
- Procedure Return
 - Use saved PC value to return to calling program
 - Restore saved values

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Encoding an Instruction Set

- How many bits for the instruction (16, 32)?
- How to allocate bits for instruction fields?
 - Op code – operation (load, add, branch)
 - Operand(s) – register/memory addresses
 - Constants (for immediate addressing)
 - Displacement (for indexing, branching)
- Fields may differ for different instructions

Outline

- A.1 Introduction
- A.2 Classifying Instruction Set Architectures
- A.3 Memory Addressing
- A.4 Type and Size of Operands
- A.5 Operations in the Instruction Set
- A.6 Instructions for Control Flow
- A.7 Encoding an Instruction Set
- A.9 MIPS Architecture

Outline

- A.9 MIPS Architecture
 - MIPS Instruction Set
 - MIPS Processor Design

The MIPS Instruction Set

- NOTE: Authors describe “RISC-V” architecture
 - We will use a specific version of this architecture: MIPS
- Why use MIPS?
 - Widely used version
 - Used in prerequisite course CS 4381/ECE 3382
 - Many resources available
 - E.g. tools that we will use in this course
- MIPS Reference Data sheet posted on Canvas

Arithmetic Operations

- Add and subtract, three operands
 - Two sources and one destination

add a, b, c # a gets b + c
- All arithmetic operations have this form

Arithmetic Example

- C code:

```
f = (g + h) - (i + j);
```

- Compiled MIPS code:

```
add t0, g, h    # temp t0 = g + h
add t1, i, j    # temp t1 = i + j
sub f, t0, t1   # f = t0 - t1
```


Register Operands

- Arithmetic instructions use register operands
- MIPS has a 32×32 -bit register file
 - Use for frequently accessed data
 - Numbered 0 to 31
 - 32-bit data called a “word”
- Assembler names
 - $\$t0, \$t1, \dots, \$t9$ for temporary values
 - $\$s0, \$s1, \dots, \$s7$ for saved variables

Register Operand Example

- C code:

$f = (g + h) - (i + j);$
– f, \dots, j in $\$s0, \dots, \$s4$

- Compiled MIPS code:

```
add $t0, $s1, $s2
add $t1, $s3, $s4
sub $s0, $t0, $t1
```

Memory Operands

- Main memory used for composite data
 - Arrays, structures, dynamic data
- To apply arithmetic operations
 - Load values from memory into registers
 - Store result from register to memory
- Memory is byte addressed
 - Each address identifies an 8-bit byte
- Words are aligned in memory
 - Address must be a multiple of 4
- MIPS is Big Endian
 - Most-significant byte at least address of a word
 - *c.f.* Little Endian: least-significant byte at least address

Memory Operand Example 1

- C code:

```
g = h + A[8];
```

– g in \$s1, h in \$s2, base address of A in \$s3

- Compiled MIPS code:

– Index 8 requires offset of 32

- 4 bytes per word

```
lw    $t0, 32($s3)    # load word
add   $s1, $s2, $t0
```

Memory Operand Example 2

- C code:

`A[12] = h + A[8];`

– `h` in `$s2`, base address of `A` in `$s3`

- Compiled MIPS code:

– Index 8 requires offset of 32

```
lw    $t0, 32($s3)    # load word
```

```
add   $t0, $s2, $t0
```

```
sw    $t0, 48($s3)    # store word
```

Registers vs. Memory

- Registers are faster to access than memory
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler must use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Immediate Operands

- Constant data specified in an instruction
`addi $s3, $s3, 4`
- No subtract immediate instruction
 - Just use a negative constant
`addi $s2, $s1, -1`

The Constant Zero

- MIPS register 0 (\$zero) is the constant 0
 - Cannot be overwritten
- Useful for common operations
 - E.g., move between registers
add \$t2, \$s1, \$zero

Unsigned Binary Integers

- Given an n-bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$

- Example

- $0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 1011_2$
 $= 0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
 $= 0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$

- Using 32 bits

- 0 to +4,294,967,295

2s-Complement Signed Integers

- Given an n-bit number

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$

- Example

- $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2$
 $= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0$
 $= -2,147,483,648 + 2,147,483,644 = -4_{10}$

- Using 32 bits

- $-2,147,483,648$ to $+2,147,483,647$

2s-Complement Signed Integers

- Bit 31 is sign bit
 - 1 for negative numbers
 - 0 for non-negative numbers
- $-(-2^{n-1})$ can't be represented
- Non-negative numbers have the same unsigned and 2s-complement representation
- Some specific numbers
 - 0: 0000 0000 ... 0000
 - -1: 1111 1111 ... 1111
 - Most-negative: 1000 0000 ... 0000
 - Most-positive: 0111 1111 ... 1111

Signed Negation

- Complement and add 1
 - Complement means $1 \rightarrow 0, 0 \rightarrow 1$

$$x + \bar{x} = 1111 \dots 111_2 = -1$$

$$\bar{x} + 1 = -x$$

■ Example: negate +2

- $+2 = 0000 \ 0000 \dots 0010_2$
- $-2 = 1111 \ 1111 \dots 1101_2 + 1$
 $= 1111 \ 1111 \dots 1110_2$

Sign Extension

- Representing a number using more bits
 - Preserve the numeric value
- In MIPS instruction set
 - `addi`: extend immediate value
 - `lb`, `lh`: extend loaded byte/halfword
 - `beq`, `bne`: extend the displacement
- Replicate the sign bit to the left
 - c.f. unsigned values: extend with 0s
- Examples: 8-bit to 16-bit
 - +2: 0000 0010 => 0000 0000 0000 0010
 - -2: 1111 1110 => 1111 1111 1111 1110