

CS/ECE 5381/7381  
Computer Architecture  
Spring 2023

Dr. Manikas


Computer Science

Lecture 21: Apr. 18, 2023

# Project 7

- Due Sat, May 6 (11:59 pm)
  - (Final Exam Period)
- Xcelium tool – more Verilog
- Assignment:
  - Develop Verilog code to implement the MIPS Data Path Control Unit
  - (common control signals – recall from App. A)

# Datapath With Control



# Project 7

- Control unit will need to handle the following MIPS instructions: **sub**, **addi**, and **lw**
- Control unit will need to output values for the following control signals:
  - **RegDst**
  - **ALUSrc**
  - **MemRead**
  - **MemWrite**
  - **MemtoReg**

# Thread-Level Parallelism

(Chapter 5, Hennessy and Patterson)

Note: some course slides adopted  
from publisher-provided material

# Outline

- 5.1 Introduction
- 5.2 Centralized Shared-Memory Architectures

# MESI Protocol – Cache Line States




- **(M) Modified** - cache line has been modified, is different from main memory - is the only cached copy. (multiprocessor 'dirty')
- **(E) Exclusive** - cache line is the same as main memory and is the only cached copy
- **(S) Shared** - Same as main memory but copies may exist in other caches.
- **(I) Invalid** - Line data is not valid (as in simple cache)

# MESI Protocol Operation

- Operation can be described informally by looking at action in local processor
  - A. Read Hit
  - B. Read Miss
  - C. Write Hit
  - D. Write Miss
- More formally by state transition diagram
- We will use a working example to illustrate

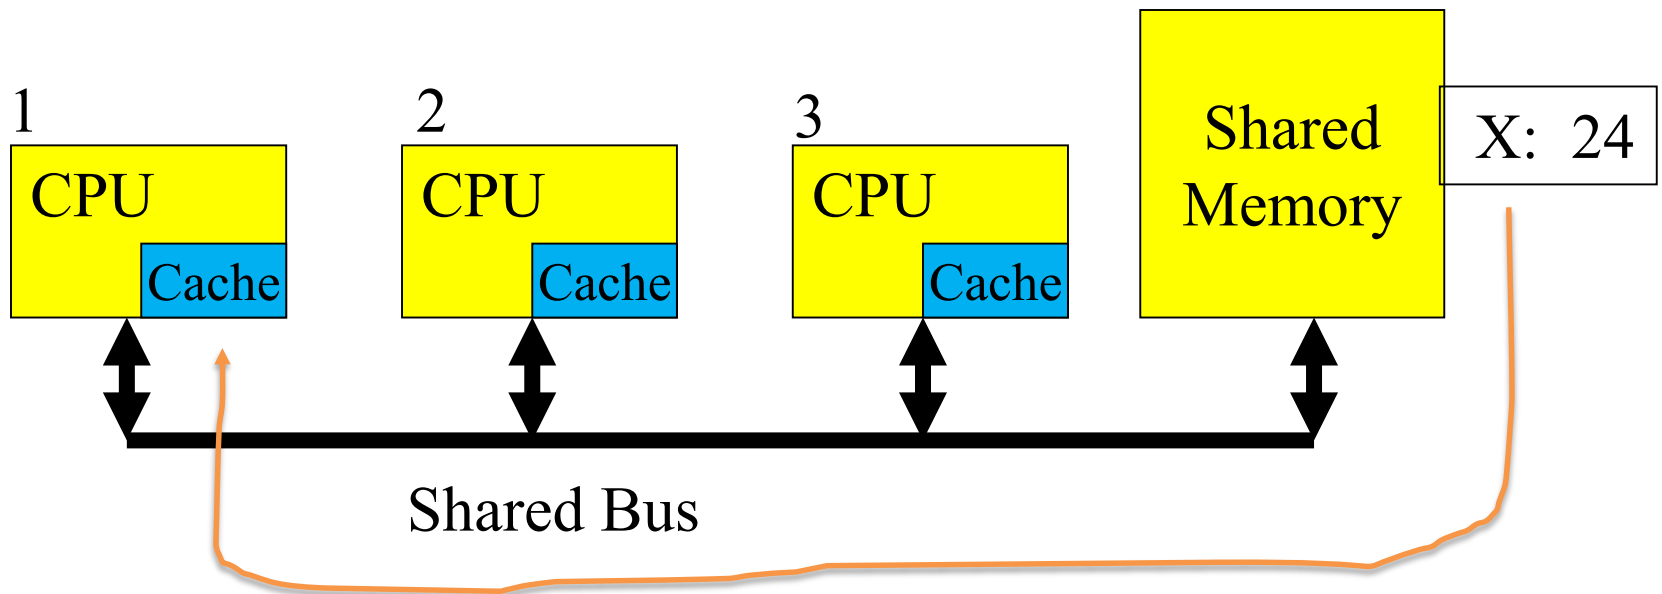


# Read Hit State Diagram




# Local Read Miss (1)

- **If no copies in other caches**
  - Processor must make a bus request to main memory
  - Value read to local cache
  - Cache line state is now E (Exclusive)
    - Cache line same as main memory and is only cached copy




Read miss for Cache 1: get X from shared (main) memory and read into appropriate cache line for cache 1 (assume X is in block 257)

### Cache 1 MESI state diagram




# Local Read Miss (2)


- **Else, if another cache has an exclusive copy of item (state E)**
  - Processor makes a bus request to main memory
  - Snooping cache puts copy of item on bus
  - Value read to local cache
  - Both cache line states are now S (Shared)
    - Both cache lines now have same value as main memory



Cache 1 MESI state diagram

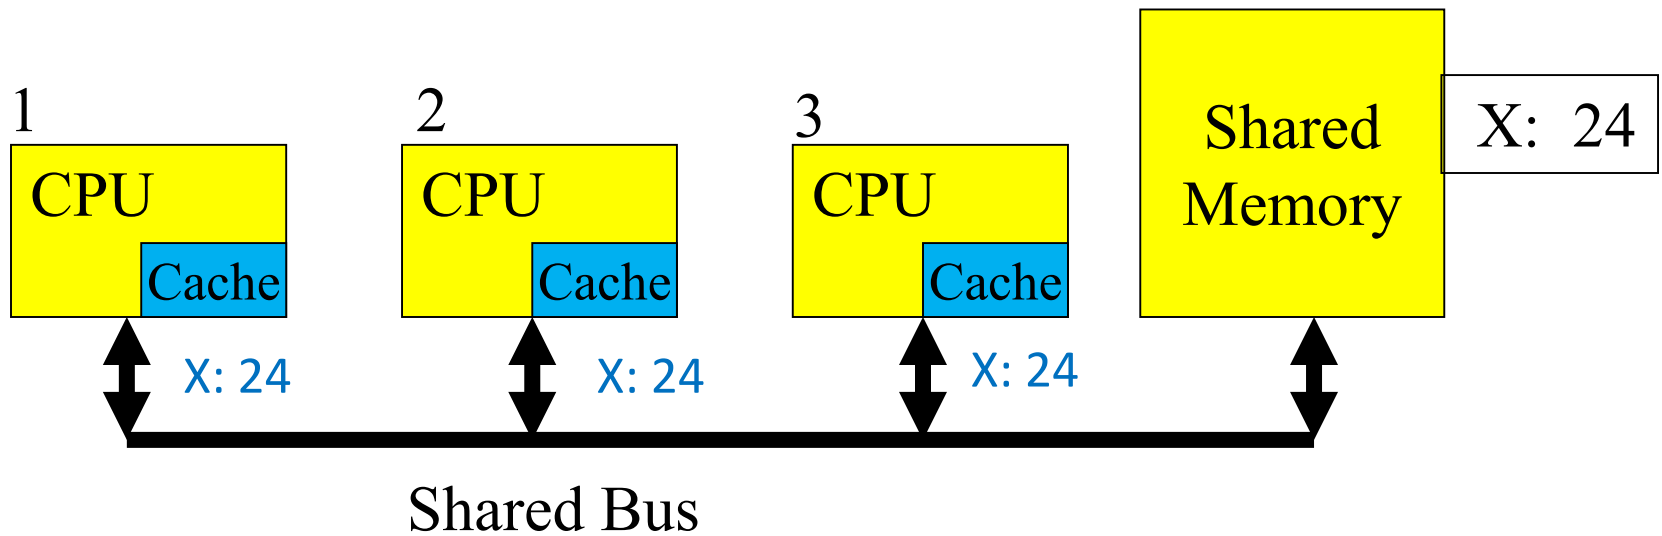


Cache 2 MESI state diagram




# Local Read Miss (3)


- **Else, if other caches have a shared copy of item (state S)**
  - Processor makes a bus request to main memory
  - **One cache** puts copy of item on bus
  - Value read to local cache
  - Local cache state is now S
  - Other caches with copies remain at state S



Cache 1 MESI state diagram



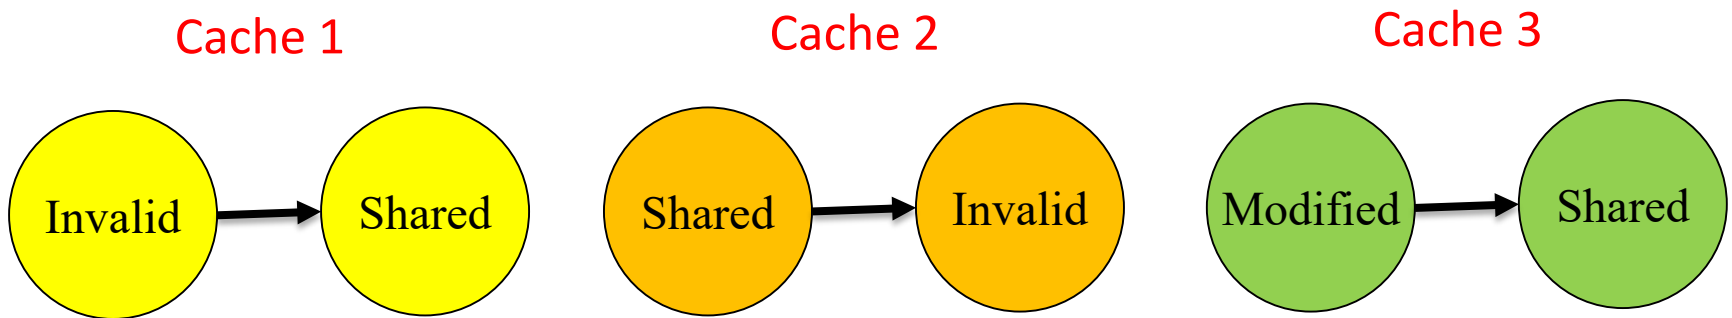

Cache 2 and 3 MESI state diagram



# Local Read Miss (4)


- **Else, if one cache has a *modified* copy of item (state M)**
  - Processor makes a bus request to main memory
  - Cache puts copy of item on bus
  - Value read to local cache
  - Local cache state is now S
  - M cache copies value to main memory
    - State changes from M to S





# Local Write Hit (M)


- **Cache line state M: cache line was modified but main memory not yet updated**
  - Cache line is exclusive and already “dirty” (old)
  - Update local cache value
  - No state change



Cache 3 has a *modified* copy of X.

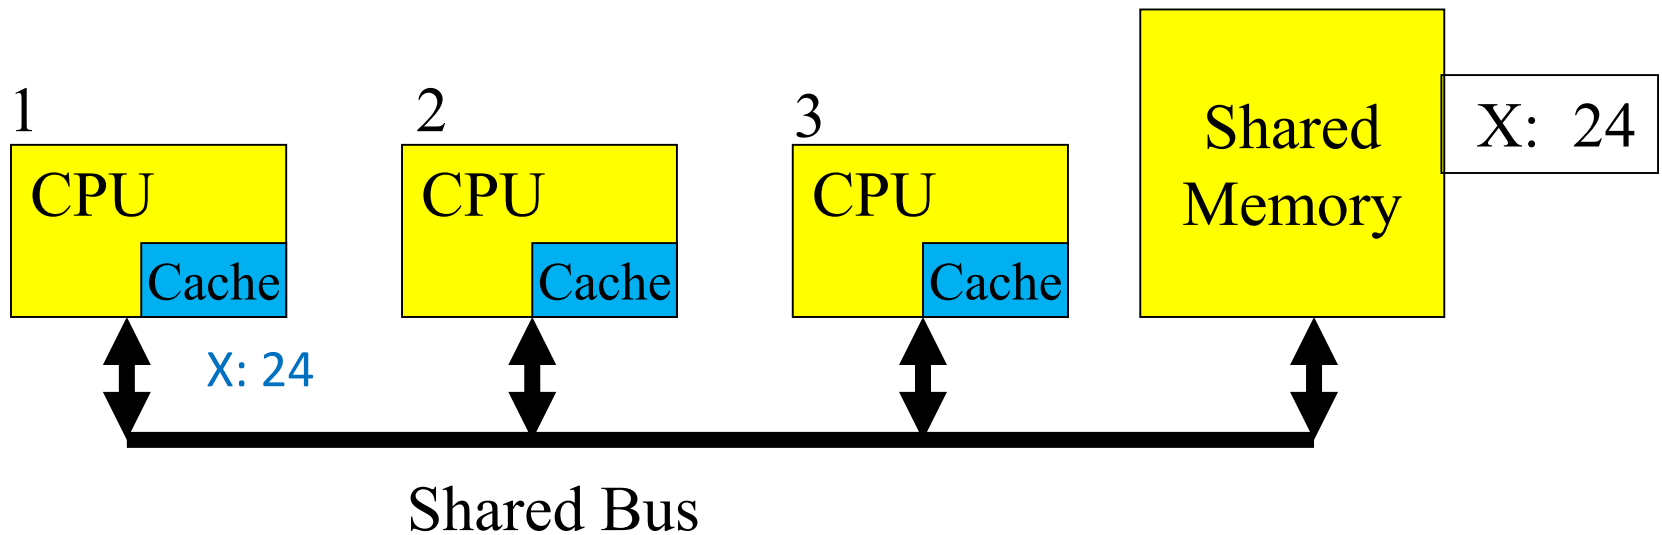
Update Cache 3 line with new value of X.

Cache 3



# Local Write Hit (E)


- **Cache line state E: cache line same as main memory and is the only cached copy**
  - Update local cache value
  - Cache line state changes to M since we are modifying value



Cache 1 line is in **Exclusive (E)** state: same value as main (shared) memory for address X



Update Cache 1 line with new value of X: **Modified (M)** state

**Cache 1**



# Local Write Hit (S)

- **Cache line state S: cache line is same as main memory and there may be copies in other caches**
  - Processor broadcasts “invalidate” on bus
  - Snooping processors with S caches change cache line state to I (Invalid)
  - Local cache value is updated
  - Local cache line state changes to M (Modified)




## D. Local Write Miss



- Write Miss: address not found in cache
  - Cache line state is I (Invalid)
- Action depends on copies of X in other caches, if they exist



# Local Write Miss (1)

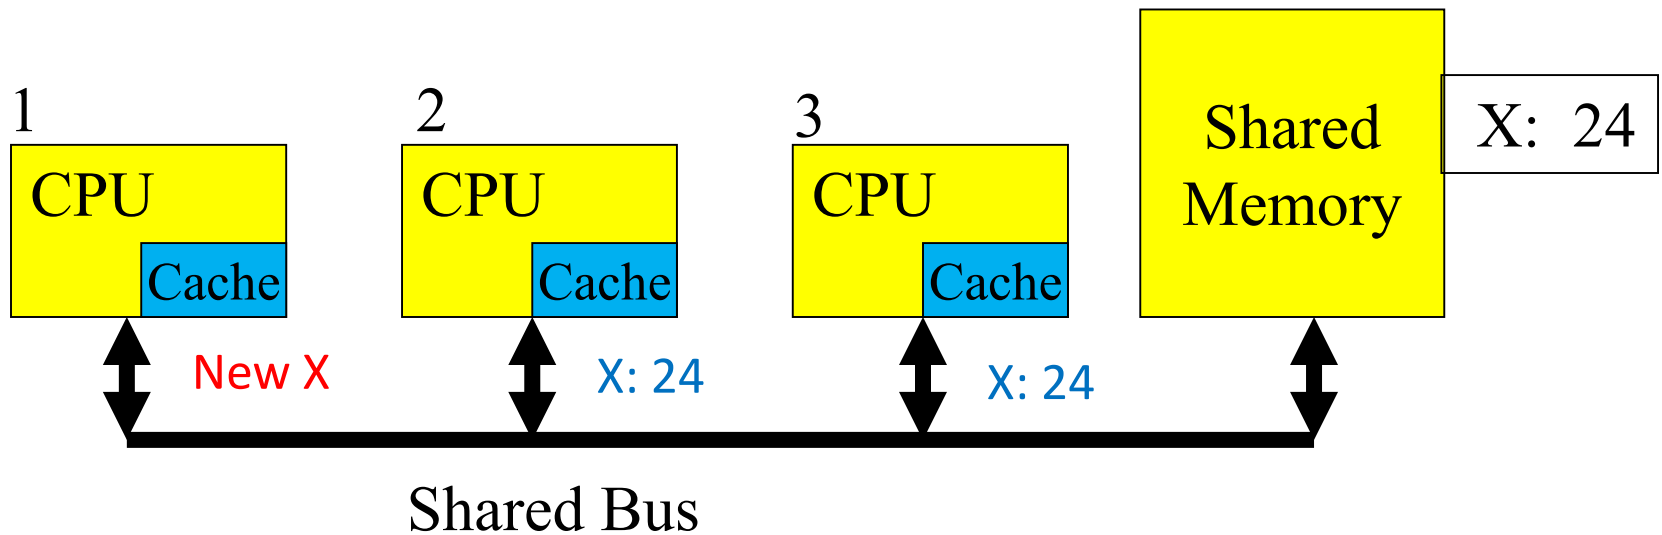
- **If no copies of X exist in other caches**
  - Update value of X in local cache line
  - Set local cache state to M (modified)





# Local Write Miss (2)

- **Else, if other copies, either one in state E or more in state S**
  - Value read from memory to local cache - bus transaction marked RWITM (read with intent to modify)
  - Snooping processors see this and set their copy state to I
  - Local copy updated & state set to M


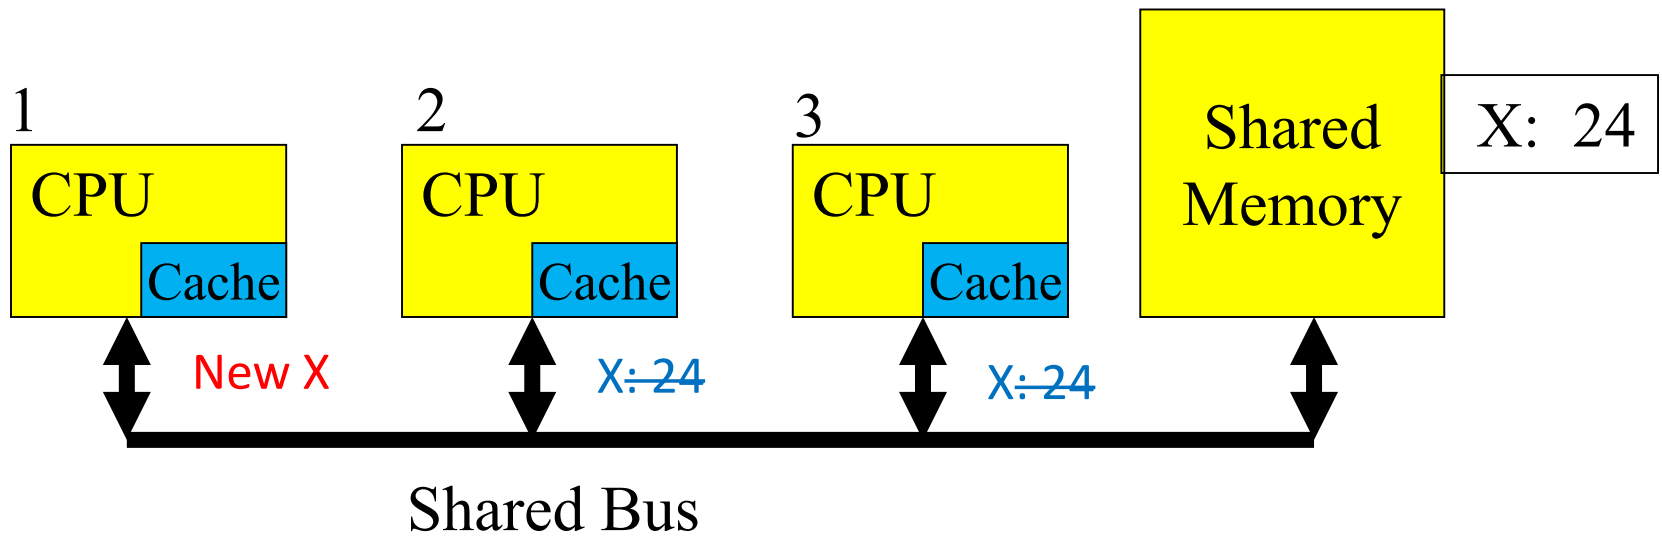


Caches 2 and 3 both have copy of current value of X

Cache 1 does NOT have copy of X, but wants to modify X. CPU 1 sends RWITM signal to bus

Caches 2 and 3 set their copies of X to “invalid”

Cache 1 modifies X



# Warehouse-Scale Computers

(Chapter 6, Hennessy and Patterson)

Note: some course slides adopted  
from publisher-provided material

# Outline

- 6.1 Introduction
- 6.2 Programming Models and Workloads
- 6.3 Architecture
- 6.4 Physical Infrastructure
- 6.5 Cloud Computing

# Internet Services Becoming Prevalent

- Internet services:
  - Webmail, search, social network, maps, video/picture sharing, storage, video streaming, language translation, ...
- The trend toward server-side or “cloud” computing
  - Moving away from PC-like clients to smaller, often mobile, devices combined with large-scaled internet services
  - Many Internet services were previously resided in the client, but now are moved to the cloud



# Internet Services Becoming Prevalent

- Characteristics of internet services
  - Multiple individual programs that interact to implement complex services to serve a massive amount of end users → one service for many

# What Computers for Internet Services?

- WSC: *warehouse-scale computer*
    - Large-scaled internet services require a computing power that can only be provided by clusters of 100s or 1000s of machines, i.e. a warehouse full of computers
    - Hardware and software resources in such a facility must work in concert as would a single computer
- **datacenter is the computer**

# Warehouse-Scale Computer (WSC)

- Large clusters – 50 to 100 thousand servers
- Housed in dedicated buildings, with special electrical and cooling requirements
- Provides Internet services
  - Google
  - Amazon



# Warehouse-Scale Computer (WSC)

- Program is an Internet service
  - May consist of tens or more individual programs that interact to implement complex end-user services such as email, search, or maps
- Software runs on hardware consisting of thousands of individual computing nodes with:
  - Networking and storage subsystems
  - Power distribution and conditioning equipment
  - Extensive cooling systems

# Warehouse-Scale Computer (WSC)

- The enclosure for these systems is in fact a building structure and often indistinguishable from a large warehouse
  - Hence the name “Warehouse”-Scale Computer

# WSC Design Factors

- Important design factors for WSC:
  - Cost-performance
    - Small savings add up
  - Energy efficiency
    - Affects power distribution and cooling
    - Work per joule
  - Dependability via redundancy
  - Network I/O
  - Interactive and batch processing workloads

# Outline

- 6.1 Introduction
- 6.2 Programming Models and Workloads
- 6.3 Architecture
- 6.4 Physical Infrastructure
- 6.5 Cloud Computing

# Programming Models and Workloads

- Batch processing framework: MapReduce
  - **Map:** applies a programmer-supplied function to each logical input record
    - Runs on thousands of computers
    - Provides new set of key-value pairs as intermediate values
  - **Reduce:** collapses values using another programmer-supplied function



# Programming Models and Workloads

- Example:
  - **map (String key, String value):**
    - `// key: document name`
    - `// value: document contents`
    - `for each word w in value`
      - `EmitIntermediate(w,"1");` `// Produce list of all words`
  - **reduce (String key, Iterator values):**
    - `// key: a word`
    - `// value: a list of counts`
    - `int result = 0;`
    - `for each v in values:`
      - `result += ParseInt(v);` `// get integer from key-value pair`
    - `Emit(AsString(result));`

# Programming Models and Workloads

- **MapReduce runtime environment schedules map and reduce task to WSC nodes**
- **Availability:**
  - Use replicas of data across different servers
  - Use relaxed consistency:
    - No need for all replicas to always agree
- **Workload demands**
  - Often vary considerably

# Outline

- 6.1 Introduction
- 6.2 Programming Models and Workloads
- 6.3 Architecture
- 6.4 Physical Infrastructure
- 6.5 Cloud Computing

# Computer Architecture of WSC

- **WSC often use a hierarchy of networks for interconnection**
- **Each 19" rack holds 48 servers connected to a rack switch**
- **Rack switches are uplinked to switch higher in hierarchy**
  - **Uplink has 6-24X times lower bandwidth**
  - **Goal is to maximize locality of communication relative to the rack**



# Storage

- **Storage options:**
  - Use disks inside the servers, or
  - Network attached storage
  - WSCs generally rely on local disks
  - Google File System (GFS) uses local disks and maintains at least three replicas

# WSC Memory Hierarchy

- Servers can access DRAM and disks on other servers

	Local	Rack	Array
DRAM latency ( $\mu$ s)	0.1	300	500
Flash latency ( $\mu$ s)	100	400	600
Disk latency ( $\mu$ s)	10,000	11,000	12,000
DRAM bandwidth (MB/s)	20,000	100	10
Flash bandwidth (MB/s)	1000	100	10
Disk bandwidth (MB/s)	200	100	10
DRAM capacity (GB)	16	1024	31,200
Flash capacity (GB)	128	20,000	600,000
Disk capacity (GB)	2000	160,000	4,800,000

# Outline

- 6.1 Introduction
- 6.2 Programming Models and Workloads
- 6.3 Architecture
- 6.4 Physical Infrastructure
- 6.5 Cloud Computing